

Temporal Order Changes a Bounded-Memory Information–Recovery Frontier

Causal list decoding and a linear rank-tail theorem

Javier Emilio Bazán Sánchez¹

¹Facultad de Ciencias, Universidad Nacional Autónoma de México, Mexico City, Mexico,
 bazan@ciencias.unam.mx

22 August 2026

Abstract

Can the temporal order of the same linear predicate change the operational power of a fixed quantum memory? A verifier streams four halves of independent EPR pairs through a device, sequestering each output before the next input. Only after commitment does a late choice request either the complete two-bit syndrome (AUDIT) or coherent recovery of all four pairs (RETURN). The device has one persistent qubit and unlimited genuinely classical transcript. Two check matrices contain the same columns: the grouped order (e_1, e_1, e_2, e_2) and the interleaved order (e_1, e_2, e_1, e_2) . The grouped order attains the full static qubit frontier and permits return fidelity $1/2$ at perfect audit. The interleaved order is instead bounded by, and attains, $1/4$. More generally, m consecutive full-rank syndrome blocks impose a temporal product law at perfect audit. The main result makes that law robust: if D is the input dimension and $\epsilon = 1 - P_A$, the spectral mass below the perfect-audit rank is at most $mD\epsilon$. Causal dimension-to-list conversion, Schur–Horn majorisation, and flagged polar recovery then give an explicit approximate-audit bound. For the interleaved instance this certifies a strict order-dependent support gap for every $3/7 < \lambda < 1$, including balanced weight. A local Pauli-completion theorem also reduces the arbitrary adaptive optimum exactly to one normalised bond-two Choi MPS. A compact three-effect MPS reaches $0.759802783851\dots$ at balanced weight and agrees with unrestricted variational searches, but the global maximum of the nonconvex MPS quotient remains uncanceled.

Keywords: bounded coherent memory; temporal order; quantum combs; entanglement recovery; syndrome decoding; Ky Fan norms.

1 Introduction

Memory dimension is a natural resource in a multitime quantum process, but dimension alone does not say how difficult a temporal task is. A memory may need to

preserve several partial functions simultaneously, and the number of live functions can change when otherwise identical inputs are reordered. This work isolates that phenomenon in a late-choice information–recovery game.

The verifier sends fresh entangled carriers one at a time. A device must commit each output before seeing the next input and does not learn the terminal task in advance. AUDIT rewards a classical syndrome of the stream, whereas RETURN rewards preservation of the unknown quantum inputs. A coherent memory can accumulate a syndrome reversibly; a classical record that helps AUDIT also damages RETURN. The question is whether a mere permutation of the syndrome columns changes the best tradeoff when the coherent memory is fixed.

It does. The pair in Fig. 1 uses the same rank-two binary code and the same one-qubit memory. Grouping equal columns realizes the complete static qubit frontier. Interleaving them creates two consecutive full-rank temporal blocks and lowers perfect-AUDIT return fidelity from $1/2$ to $1/4$. The effect is not limited to an endpoint: a robust rank-tail theorem proves a strict support-function separation throughout $3/7 < \lambda < 1$.

The remaining interior optimisation can be stated without an unbounded adaptive tree. Every normalised bond-two Choi leaf admits a local four-outcome Pauli completion whose transcripts are covariance-related copies of that leaf. Consequently the exact support is a compact Choi-MPS maximum. This reduction turns local completeness from an obstruction into a theorem; it does not make the resulting nonconvex maximum analytic automatically.

Main consequence. For the interleaved stream, every admissible strategy obeys $\beta_{H_1,2}(\lambda) \leq U_1(\lambda)$, and $U_1(\lambda) < B_{4,2}(\lambda)$ whenever $3/7 < \lambda < 1$. Thus temporal order changes the achievable information–recovery region in the interior, not only at perfect AUDIT.

The broad ingredients are established. Quantum combs and process tensors describe multitime strategies and their internal memories [1, 2, 3]. Bounded




Figure 1: The same four columns in two temporal orders. The grouped stream has one complete rank-two block. The interleaved stream has two; a one-qubit bond must cross both full-rank boundaries. AUDIT sees only the classical transcript and terminal qubit, whereas RETURN regains all committed carriers.

coherent memory has been studied through constrained-separability hierarchies and autonomous process discrimination [4, 5]. Postmeasurement-information games allow unlimited classical data beside a bounded quantum register [6, 7], while information-disturbance theory relates measurement records to later recovery [8, 9, 10]. Classical trellis complexity is famously order dependent [11, 12, 13]. The claim here is narrower: an exact same-code temporal-order separation in a common-prefix syndrome-AUDIT/all-carrier-RETURN game, together with a causal-list rank-tail bound. A focused search through 21 August 2026 located no direct antecedent of that conjunction; this is a dated collision assessment through 22 August 2026, not proof of priority.

2 Late-choice streaming model

At slot $i \in \{1, \dots, n\}$, the verifier prepares

$$|\Phi^+\rangle_{R_i A_i} = \frac{|00\rangle + |11\rangle}{\sqrt{2}}, \quad (1)$$

sends A_i to the device, receives B_i , and immediately sequesters B_i . The device applies an arbitrary adaptive instrument to A_i , fresh local ancillas, and a persistent memory M_{i-1} . Only M_i , with declared dimension at most d_i , may remain coherent across the next cut. An unrestricted finite classical transcript is free, but it must be genuinely dephased: retaining its coherent purification would be an additional quantum memory.

After slot n , a hidden choice selects one task. In AUDIT, the references are measured in the computational basis, producing a uniform word $X \in \mathbb{F}_q^n$. The device receives neither references nor sequestered carriers and must output the full syndrome $HX \in \mathbb{F}_q^r$ using only its transcript and terminal memory. Its unconditional success is P_A . In RETURN, a transcript-conditioned decoder acts jointly on all B_i and the terminal memory, without reference access. The verifier tests all EPR pairs and the charged memory reset. Its unconditional acceptance probability is F_R .

For $0 \leq \lambda \leq 1$, define the support function

$$\beta_{H,d}^{\text{stream}}(\lambda) = \sup_{\mathcal{S}} [\lambda P_A(\mathcal{S}) + (1 - \lambda) F_R(\mathcal{S})], \quad (2)$$

where the supremum includes arbitrary finite-outcome adaptive non-QND strategies satisfying the interface.

No postselection is allowed. Immediate sequestration matters: past carriers cannot silently enlarge a later memory.

For the canonical binary pair,

$$H_G = \begin{pmatrix} 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 \end{pmatrix}, \quad H_I = \begin{pmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \end{pmatrix}. \quad (3)$$

They differ only by a coordinate permutation. Both have four equiprobable syndromes, and the persistent memory is one qubit.

3 Static ceiling and grouped attainment

Ignoring streaming order gives a useful terminal-dimension relaxation. For a four-label ensemble and a terminal qubit, the exact support is

$$B_{4,2}(\lambda) = \frac{1 + \sqrt{\lambda^2 + (1 - \lambda)^2}}{2}. \quad (4)$$

The bound follows by refining the classical flag, applying the dimension-two minimum-error bound to AUDIT, and the flagged recovery inequality to RETURN. For a normalized four-component likelihood row, the extremal profile has two equal large entries and two equal small entries. The remaining optimization is the Euclidean norm in (4).

Proposition 3.1 (Grouped frontier). *For H_G and a persistent qubit,*

$$\beta_{H_G,2}^{\text{stream}}(\lambda) = B_{4,2}(\lambda) \quad (0 \leq \lambda \leq 1). \quad (5)$$

An exposed parametrization is

$$P_A = \frac{1+t}{2}, \quad F_R = \frac{1 + \sqrt{1-t^2}}{2}, \quad 0 \leq t \leq 1. \quad (6)$$

The device coherently accumulates the first parity, weakly records it with strength t , uncomputes the carrier-dependent part, and hands the same qubit to the second parity. The classical flag and terminal qubit identify the full syndrome in AUDIT; the carrier outputs undo the coherent updates in RETURN. At perfect AUDIT, (6) gives $F_R = 1/2$.

4 Temporal fragmentation at perfect audit

The endpoint extends well beyond the four-slot example. Let $H \in \mathbb{F}_q^{r \times n}$ be split into m consecutive blocks,

$$H = [H^{(1)} | \dots | H^{(m)}], \quad \text{rank } H^{(j)} = r. \quad (7)$$

Put $N = q^r$. Let at most $d_j \leq N$ coherent dimensions cross the boundary after block j , including the terminal boundary. The transcript remains free and classical.

Theorem 4.1 (Perfect-AUDIT temporal product law). *Under (7),*

$$P_A = 1 \implies F_R \leq \prod_{j=1}^m \frac{d_j}{N}. \quad (8)$$

For uniform $d = q^s$, the bound is attained on $H = [I_r | \dots | I_r]$ by retaining s syndrome coordinates coherently and recording the other $r - s$ coordinates in each block.

Each full-rank block introduces an independently uniform syndrome contribution. Perfect terminal decoding constrains every refined Kraus leaf to at most d_j cumulative labels at boundary j . Intersecting those constraints leaves a support fraction $\prod_j (d_j/N)$. Optimal flagged recovery converts that rank fraction into (8).

For H_G , one can pack only one consecutive full-rank block, so the endpoint is $1/2$. For $H_I = [I_2 | I_2]$, there are two and the endpoint is $1/4$. Both are attained. Repeating each basis column m times gives a larger same-columns separation: batching equal columns yields d/q^r , while cycling through the basis m times yields $(d/q^r)^m$. The ratio grows as $(q^r/d)^{m-1}$. Consecutive trellis sectionalization itself is classical prior art [14]; the new statement is the late-choice recovery law attached to it.

5 Causal lists and a linear spectral tail

Perfect decoding forces exact low rank. Approximate decoding requires a quantitative substitute. Let $D = q^n$ and

$$\alpha = \prod_{j=1}^m \frac{d_j}{N}, \quad k = \alpha D = q^{n-mr} \prod_j d_j. \quad (9)$$

Refine every instrument outcome to a single Kraus label and reveal the label to both late decoders. This only enlarges the strategy class. A complete transcript c has a sequential Kraus operator K_c and effect $G_c = K_c^\dagger K_c$. Let

$$t_c = \sum_{\ell > k} \lambda_\ell(G_c), \quad (10)$$

where eigenvalues are nonincreasing.

Theorem 5.1 (Linear rank-tail robustness). *For every admissible strategy under (7), with $\epsilon = 1 - P_A$,*

$$\sum_c t_c \leq m D \epsilon. \quad (11)$$

Consequently, for $\theta = \min\{m\epsilon, 1 - \alpha\}$,

$$F_R \leq \alpha + (1 - 2\alpha)\theta + 2\sqrt{\alpha(1 - \alpha)\theta(1 - \theta)}. \quad (12)$$

The proof uses a dimension-to-list conversion. At boundary j , pull the terminal AUDIT measurement backward through the future comb, average over the independent future inputs, and relabel the final syndrome as the cumulative prefix syndrome. This produces a valid POVM on the d_j -dimensional cut memory with the same global success P_A . A dimension- d_j POVM implies a classical list L_j of at most d_j labels whose event probability is at least P_A .

For a complete transcript, intersect the m causal list events. Their failure probability is at most $m\epsilon$. The cumulative block labels are in bijection with the individual block syndromes, and full rank makes every fibre uniform. The intersection therefore contains at most k computational words. Schur–Horn majorisation turns its coordinate mass into the top- k spectral mass of G_c , proving (11). This route avoids the extra square-root losses of approximate-injectivity arguments.

To obtain RETURN, split the singular values above and below rank k :

$$\text{Tr } \sqrt{G_c} \leq \sqrt{k(m_c - t_c)} + \sqrt{(D - k)t_c}, \quad m_c = \text{Tr } G_c. \quad (13)$$

Instrument completeness gives $\sum_c m_c = D$. Cauchy–Schwarz sums (13), while Ky Fan anti-norm superadditivity gives the universal tail ceiling $\sum_c t_c \leq D - k$ [15]. Substitution yields (12). The right-hand side is a squared overlap of two real unit vectors, so its monotonicity on the allowed interval is immediate.

6 An explicit interior order gap

For H_I , $m = 2$ and $\alpha = 1/4$. Hence, with

$$\theta = \min\{2(1 - P_A), 3/4\}, \quad (14)$$

every strategy satisfies

$$F_R \leq \frac{1}{4} + \frac{1}{2}\theta + \frac{\sqrt{3}}{2}\sqrt{\theta(1 - \theta)}. \quad (15)$$

Near perfect AUDIT this becomes

$$F_R \leq \frac{1}{4} + (1 - P_A) + \sqrt{\frac{3}{2}(1 - P_A)(2P_A - 1)}. \quad (16)$$

The square-root exponent is optimal: a legal weak-record family has a $\sqrt{1 - P_A}$ correction. The leading constant is not claimed sharp.

Maximizing (15) in the support direction gives a closed certificate.

Corollary 6.1 (Certified interior separation). *For every $0 \leq \lambda \leq 1$,*

$$\beta_{H_I, 2}^{\text{stream}}(\lambda) \leq U_I(\lambda) = \frac{1}{2} + \frac{\lambda}{4} + \frac{1}{4}\sqrt{7\lambda^2 - 10\lambda + 4}. \quad (17)$$

Moreover,

$$U_I(\lambda) < B_{4,2}(\lambda) \quad \text{for } \frac{3}{7} < \lambda < 1. \quad (18)$$

Set $\theta = \sin^2 \phi$, $0 \leq \phi \leq \pi/3$. The support becomes a positive 2×2 quadratic form in $(\cos \phi, \sin \phi)$; its largest eigenvalue is (17). Eliminating the two radicals when comparing (17) with (4) gives $-4\lambda^2(\lambda-1)(7\lambda-3)$, with the unsquared signs fixing the interval in (18).

At balanced weight,

$$\beta_{H_1,2}^{\text{stream}}(1/2) \leq \frac{5}{8} + \frac{\sqrt{3}}{8} = 0.841506350946\dots, \quad (19)$$

whereas the grouped value is $B_{4,2}(1/2) = 0.853553390593\dots$. This proves an interior gap without assuming QND, Clifford, projective, covariant, or binary-outcome instruments.

7 Exact leaf reduction and compact construction

Write a fine transcript leaf as $K : A_1 \cdots A_4 \rightarrow B_1 \cdots B_4 M$. Its vectorisation is an open-boundary MPS whose three temporal bond ranks are at most two. For a normalised leaf, define

$$A_H(K) = \max_{\{E_s\}} \sum_x \text{Tr}[E_{Hx} \text{Tr}_{B^4} K|x\rangle\langle x|K^\dagger], \quad (20)$$

$$R(K) = \frac{1}{16} \left(\text{Tr} \sqrt{K^\dagger K} \right)^2, \quad \text{Tr} K^\dagger K = 1. \quad (21)$$

Theorem 7.1 (Homogeneous Choi-MPS reduction). *For the interleaved game,*

$$\beta_{H_1,2}^{\text{stream}}(\lambda) = \max_{\substack{\text{Tr} K^\dagger K = 1 \\ \text{TT-rank}_i(|K\rangle) \leq 2}} [\lambda A_{H_1}(K) + (1-\lambda)R(K)]. \quad (22)$$

Four classical outcomes per slot suffice.

The upper direction is a homogeneous averaging argument: every complete strategy is a probability distribution over its normalised fine leaves, and both branch scores average with the same leaf weights. For the converse, put an arbitrary feasible leaf in row-isometric MPS gauge, with local tensors T_i satisfying

$$\text{Tr}_{A_i}(T_i^\dagger T_i) = \mathbb{I}_{M_{i-1}}. \quad (23)$$

For each Pauli $P \in \{I, X, Y, Z\}$ use the local Kraus map $T_i(P \otimes \mathbb{I})/\sqrt{2}$. The Pauli one-design identity makes these four maps complete. Every four-slot transcript is K right-multiplied by a tensor product of Paulis and scaled by $1/4$. RETURN is invariant, while the known Pauli bit flips merely relabel the linear syndrome in AUDIT. Thus all 256 leaves have the same homogeneous score, proving (22).

The strongest compact construction found inside (22) has three nonzero rank-one effects. It admits a two-variable support formula. Put

$$\begin{aligned} g &= 1 - t/2, & z &= -t/(2-t), \\ e &= t(1+r)/2, \\ \Delta &= \sqrt{(1-z^2)(1-r^2)}, \\ h &= \frac{g}{2}(1+zr+\Delta), \\ c_0 &= \sqrt{t} + \sqrt{1-t}, & c_1 &= \sqrt{e} + \sqrt{1-e}, \end{aligned} \quad (24)$$

and let $c_\emptyset = \sqrt{2}$ for $t \geq 1/2$, while $c_\emptyset = c_0$ for $t < 1/2$. Define

$$\begin{aligned} \mathbf{c} &= (c_0, c_\emptyset, \sqrt{2}c_1)^\top, \\ M_\lambda(t, r) &= \lambda \text{diag}(t, 0, h) + \frac{1-\lambda}{8} \mathbf{c} \mathbf{c}^\top. \end{aligned} \quad (25)$$

Then the physical lower bound is

$$L_I(\lambda) = \max \left\{ 1 - \frac{\lambda}{2}, \max_{\substack{0 \leq t \leq 1 \\ -1 \leq r \leq 1}} \lambda_{\max}(M_\lambda(t, r)) \right\}. \quad (26)$$

The Perron eigenvector fixes the four middle-cut priors. The first two slots prepare those four pure states; slot three implements the three-effect POVM; slot four generates a uniform bit and XORs the binary decision into the terminal qubit. Orthogonal emitted states make $K^\dagger K$ diagonal, so the polar RETURN value in (21) is attained exactly.

At balanced weight, (26) gives

$$\begin{aligned} t &= 0.458073977\dots, \\ r &= 0.016373524\dots, \\ P_A &= 0.620085075586\dots, \\ F_R &= 0.899520492117\dots, \\ L_I(1/2) &= 0.759802783851444\dots \end{aligned} \quad (27)$$

It crosses the no-record line at $\lambda_c = 0.441437845098\dots$. An independent NumPy construction checks all temporal ranks, local row isometries, Pauli completeness, and direct scores; its largest residual is below 9×10^{-16} .

Unrestricted complex Choi-MPS continuation, eight multistarts of the larger pinched cq-instrument relaxation, and a relaxation of the first two slots all agree with (26) on the tested grid. This is substantially stronger evidence than the earlier ternary tree value $0.759448970317\dots$, and Pauli completion proves that the compact leaf is itself physical. It is still not a global upper certificate for the nonconvex maximum in (22); equality between L_I and the exact support remains a conjecture.

8 Relation to neighbouring disciplines

The result sits at a specific intersection. Bounded-storage cryptography and postmeasurement information already turn a quantum dimension into a late decoding limitation [6, 7]. Quantum-comb memory cost already separates persistent coherent dimensions from classical assistance [3, 4]. Sequential state generation identifies persistent ancilla dimension with a temporal tensor-network bond [16]. Trellises, sectionalization, and matroid pathwidth already quantify order-sensitive live constraints in linear codes [11, 14, 13]. Nondestructive discrimination connects classical labels with recoverable entanglement in a spatial LOCC setting [17, 18].

The present bridge is operational rather than terminological. A causal dimension bound is converted into one




Figure 2: Support values for the same-code pair. The blue curve is the exact grouped frontier. The orange curve is the rigorous interleaved upper certificate; the shaded region is excluded solely by temporal order for $3/7 < \lambda < 1$. The dashed curve is the compact three-effect MPS lower bound; the hollow diamond is the earlier ternary strategy that it supersedes. The unclosed vertical distance is not an error bar: it is the remaining global MPS-maximisation problem.

classical list at every full-rank temporal section. The intersection size controls a Ky Fan tail of each instrument leaf, which in turn controls all-carrier entanglement recovery. None of the nearby papers located in the focused review states this chain for the same-code, late-choice syndrome/return interface. Conversely, the work does not claim that order dependence, trellis width, bounded quantum memory, or information–disturbance is new in isolation.

9 Limits and open exact frontier

The theorem is conditional on a trusted resource boundary. Every coherent system crossing a block cut must count toward d_j ; the classical transcript must be dephased; past carriers must be unavailable in AUDIT; and RETURN must include all failures. Relaxing any of these assumptions changes the task. Visible memory reset does not prove that an inaccessible environment contains no record.

The interval (18) resolves the existence of an interior order effect, not the full interleaved Pareto boundary. At balanced weight the best verified lower strategy and rigorous upper certificate are

$$\begin{aligned} 0.759802783851\dots &\leq \beta_{H_1,2}^{\text{stream}}(1/2) \\ &\leq 0.841506350946\dots \end{aligned} \quad (28)$$

The Choi-MPS reduction proves that local completeness and unbounded adaptive outcome trees are no longer the obstruction. Closing the gap now requires either a better bond-two MPS leaf or a global upper certificate for the nonconvex quotient in (22). The optimal

coefficient in the near-endpoint square-root law and extensions to partially ranked block sequences are also open.

Nothing in the observed statistics selects an interpretation of quantum mechanics. The result is a standard operational statement about coherent memory, temporal order, and recoverability. It neither observes alternative worlds nor communicates with an environmentally decohered branch.

10 Conclusion

Permuting the columns of one linear syndrome can change the attainable information–recovery region of a fixed coherent memory. Grouping equal columns allows a one-qubit device to attain the static frontier; interleaving them creates an additional full-rank temporal boundary, squares the perfect-AUDIT recovery penalty, and produces a rigorously certified interior gap. Causal list decoding provides the robust mechanism: approximate terminal knowledge yields small lists at every memory cut, their intersection bounds a spectral tail, and that tail limits coherent recovery. This narrow theorem, rather than a many-worlds narrative, is the contribution that survives the originality audit. Local Pauli completion further removes the arbitrary adaptive tree exactly and leaves one finite bond-two Choi-MPS maximum. The three-effect construction gives the strongest verified lower curve, but the remaining global maximum is deliberately left open.

A Dimension-to-list lemma

Lemma A.1. *Let $\{\rho_x\}_{x \in \mathcal{X}}$ be subnormalised states on a d -dimensional system and $\{E_x\}_{x \in \mathcal{X}}$ a POVM. If $p_x = \text{Tr } \rho_x$, then*

$$\sum_x \text{Tr}(E_x \rho_x) \leq \max_{L \subseteq \mathcal{X}, |L| \leq d} \sum_{x \in L} p_x. \quad (29)$$

Proof. Positivity gives $\text{Tr}(E_x \rho_x) \leq p_x \|E_x\|_\infty$. The weights $w_x = \|E_x\|_\infty$ satisfy $0 \leq w_x \leq 1$ and

$$\sum_x w_x \leq \sum_x \text{Tr } E_x = d. \quad (30)$$

The resulting fractional d -unit linear program assigns unit weight to the d largest p_x , proving (29). $\square$

B Pulled-back boundary measurements

Write a word as blocks $x = (x_1, \dots, x_m)$ and define

$$v_j = H^{(j)} x_j, \quad u_j = v_1 + \dots + v_j. \quad (31)$$

Let a_j be the refined transcript through boundary j . Conditioned on a_j and $u_j = u$, trace all sequestered prefix carriers and within-fibre variables, producing a subnormalised state $\rho_{a_j, u}^{(j)}$ on the cut memory.

Fix the future block inputs and compose the suffix comb with the terminal AUDIT POVM. Relabel its final answer u_m as

$$u_j = u_m - (v_{j+1} + \dots + v_m). \quad (32)$$

Averaging these effects over independent uniform future inputs gives a POVM on the cut memory. Summed over prefix transcripts, its success is the original P_A . The construction remains valid when the cut memory is entangled with past carriers: the future comb acts only on the memory and the carriers are traced in AUDIT.

Applying Appendix A at every prefix transcript gives a classical list $L_j(a_j) \subseteq \mathbb{F}_q^r$, $|L_j(a_j)| \leq d_j$, such that

$$\Pr[U_j \in L_j(A_j)] \geq P_A. \quad (33)$$

The POVMs for different j need not be jointly measurable; the lists are proof objects, not simultaneous physical measurements.

For a complete transcript c , define

$$S_c = \{x : u_j(x) \in L_j(a_j(c)) \text{ for all } j\}. \quad (34)$$

The map $(u_1, \dots, u_m) \leftrightarrow (v_1, \dots, v_m)$ is bijective. Every full-rank block has fibre size $q^{n_j - r}$, hence

$$|S_c| \leq q^{n - mr} \prod_j d_j = k. \quad (35)$$

The union bound applied to (33) yields

$$\sum_c \sum_{x \notin S_c} \|K_c|x\rangle\|^2 \leq mD(1 - P_A). \quad (36)$$

The diagonal of G_c is precisely $\|K_c|x\rangle\|^2$. Schur–Horn majorisation implies that the diagonal mass on any set of at most k coordinates is no greater than the sum of the top k eigenvalues. Subtracting from $\text{Tr } G_c$ in (36) proves (11).

C From spectral tail to recovery

For a fine-grained leaf, optimal polar correction contributes at most

$$F_{R,c} \leq \frac{(\text{Tr } \sqrt{G_c})^2}{D^2}. \quad (37)$$

This already allows a joint decoder on every emitted carrier and terminal memory; dropping the reset projector is an upper relaxation. Applying Cauchy–Schwarz separately to the first k and last $D - k$ singular values gives (13). Let $T = \sum_c t_c = \theta' D$. Since $\sum_c G_c = \mathbb{I}_D$, one has $\sum_c m_c = D$. Summing the square of (13) and applying Cauchy–Schwarz to $\sum_c \sqrt{t_c(m_c - t_c)}$ gives

$$F_R \leq \alpha + (1 - 2\alpha)\theta' + 2\sqrt{\alpha(1 - \alpha)\theta'(1 - \theta')}. \quad (38)$$

The sum of the smallest $D - k$ eigenvalues is a homogeneous concave spectral functional, hence superadditive. Therefore

$$T \leq \sum_{\ell > k} \lambda_\ell \left(\sum_c G_c \right) = D - k. \quad (39)$$

Together with (11), this gives $0 \leq \theta' \leq \min\{m(1 - P_A), 1 - \alpha\}$. The recovery expression is increasing throughout that interval, proving (12).

D Homogeneous Choi-MPS completion

This appendix proves (22). Refine every classical transcript of a physical four-slot strategy until it has one global Kraus operator K_c , and put

$$G_c = K_c^\dagger K_c, \quad t_c = \text{Tr } G_c. \quad (40)$$

Instrument completeness gives $\sum_c G_c = \mathbb{I}_{16}$ and $\sum_c t_c = 16$. For $t_c > 0$, write $\hat{K}_c = K_c / \sqrt{t_c}$. Because the terminal measurement and flagged recovery may depend on c , the two branch scores decompose exactly as

$$P_A = \sum_c \frac{t_c}{16} A_{H_i}(\hat{K}_c), \quad F_R = \sum_c \frac{t_c}{16} R(\hat{K}_c). \quad (41)$$

The weights in (41) sum to one. Every leaf has temporal TT ranks at most two, so its support score is no larger than the right-hand side of (22). This proves the upper direction independently of the number or adaptivity of the outcomes.

For the converse, take any normalised operator K admitted by the maximum in (22). Successive Schmidt decompositions put its vectorisation in row-isometric MPS gauge. Its local tensors

$$T_i : A_i \otimes M_{i-1} \longrightarrow B_i \otimes M_i \quad (42)$$

obey

$$\text{Tr}_{A_i}(T_i^\dagger T_i) = \mathbb{I}_{M_{i-1}}. \quad (43)$$

Zero Schmidt subspaces may be removed and smaller bonds padded, so this gauge does not increase the allowed memory dimension.

For $P \in \{I, X, Y, Z\}$, define the four local Kraus maps

$$L_{i,P} = 2^{-1/2} T_i(P \otimes \mathbb{I}_{M_{i-1}}). \quad (44)$$

The Pauli one-design identity

$$\sum_P (P^\dagger \otimes \mathbb{I}) Y (P \otimes \mathbb{I}) = 2 \mathbb{I}_{A_i} \otimes \text{Tr}_{A_i} Y \quad (45)$$

together with (43) yields $\sum_P L_{i,P}^\dagger L_{i,P} = \mathbb{I}_{A_i M_{i-1}}$. Thus (44) is a complete causal instrument at every slot and its outcome may be dephased immediately.

For a full Pauli transcript $\mathbf{P} = (P_1, \dots, P_4)$, contraction gives

$$K_{\mathbf{P}} = \frac{1}{4} K(P_1 \otimes \dots \otimes P_4). \quad (46)$$

All 256 leaves have $\text{Tr } K_{\mathbf{P}}^\dagger K_{\mathbf{P}} = 1/16$ and hence probability $1/256$ on the maximally mixed input. Right-unitary invariance preserves $R(K)$. If $P_i = X^{a_i} Z^{b_i}$, then $P_i|x_i\rangle$ is a phase times $|x_i + a_i\rangle$; subtracting the known shift $H_1(a_1, \dots, a_4)$ from the terminal answer preserves $A_{H_1}(K)$. Every conditional leaf therefore has the same two scores as K . Their complete mixture attains its homogeneous quotient and proves the lower direction of (22).

E Reproducibility and claim boundary

The code, stored instrument, tests, literature ledger, and both manuscript sources are public at <https://github.com/asim-v/carmenq>. The commands

```
python -m pytest -q
python scripts/verify_interleaved_counterexample.py
python scripts/verify_compact_interleaved_candidate.py
python scripts/generate_order_gap_figure.py
python scripts/build_order_pdf.py
```

reproduce the software checks, both balanced constructions, plotted bounds, and this PDF. Numerical searches are explicitly lower-bound generators. The proved claims are (8), (11), (12), (22), the attainability of (26), and the interval (18). Equality between (26) and the exact interleaved interior curve is not claimed.

References

- [1] Giulio Chiribella, Giacomo Mauro D’Ariano, and Paolo Perinotti. Theoretical framework for quantum networks. *Physical Review A*, 80:022339, 2009.
- [2] Felix A. Pollock, César Rodríguez-Rosario, Thomas Frauenheim, Mauro Paternostro, and Kavan Modi. Non-markovian quantum processes: Complete framework and efficient characterization. *Physical Review A*, 97:012127, 2018.
- [3] Alessandro Bisio, Giacomo Mauro D’Ariano, Paolo Perinotti, and Michal Sedlák. Memory cost of quantum protocols. *Physical Review A*, 85:032333, 2012.
- [4] Ties-Albrecht Ohst, Shijun Zhang, Hai Chau Nguyen, Martin Plávala, and Marco Túlio Quintino. Characterising memory in quantum channel discrimination via constrained separability problems. *Quantum*, 10:1988, 2026.
- [5] Magdalini Zonnios and Felix C. Binder. Distinguishing quantum processes with bounded coherent memory, June 2026. Preprint; not peer reviewed at the review cutoff.
- [6] Robert König, Ueli Maurer, and Renato Renner. On the power of quantum memory. *IEEE Transactions on Information Theory*, 51(7):2391–2401, 2005.
- [7] Manuel A. Ballester, Stephanie Wehner, and Andreas Winter. State discrimination with post-measurement information. *IEEE Transactions on Information Theory*, 54(9):4183–4198, 2008.
- [8] Konrad Banaszek. Fidelity balance in quantum operations. *Physical Review Letters*, 86:1366–1369, 2001.
- [9] M. Gregoratti and Reinhard F. Werner. Quantum lost and found. *Journal of Modern Optics*, 50(6–7):915–933, 2003.
- [10] Chung-Yun Hsieh, Armin Tavakoli, Huan-Yu Ku, and Paul Skrzypczyk. Resource theory of interactive quantum instruments. *Physical Review A*, 113:062445, 2026.
- [11] G. David Forney, Jr. Dimension/length profiles and trellis complexity of linear block codes. *IEEE Transactions on Information Theory*, 40(6):1741–1752, 1994.
- [12] Robert J. McEliece. On the BCJR trellis for linear block codes. *IEEE Transactions on Information Theory*, 42(4):1072–1092, 1996.
- [13] Navin Kashyap. Matroid pathwidth and code trellis complexity. *SIAM Journal on Discrete Mathematics*, 22(1):256–272, 2008.
- [14] Alec Lafourcade and Alexander Vardy. Optimal section-alization of a trellis. *IEEE Transactions on Information Theory*, 42(3):689–703, 1996.
- [15] Rajendra Bhatia. *Matrix Analysis*, volume 169 of *Graduate Texts in Mathematics*. Springer, New York, 1997.
- [16] Bikun Li, Sophia E. Economou, and Edwin Barnes. Photonic resource state generation from a minimal number of quantum emitters. *npj Quantum Information*, 8:11, 2022.
- [17] Bohdan Bilash, Youngrong Lim, Hyukjoon Kwon, Yosep Kim, Hyang-Tag Lim, Wooyeong Song, and Yong-Su Kim. Nondestructive discrimination of bell states between distant parties. *Physical Review A*, 110:042407, 2024.
- [18] Youngrong Lim, Minki Hhan, and Hyukjoon Kwon. Trade-off between information gain and disturbance in local discrimination of entangled quantum states. *Quantum Science and Technology*, 10(2):025048, 2025.